

SANSKRIT DOCUMENT

Retrieval-Augmented Generation (RAG) System

Technical Report

Objective

Design and implement a CPU-based Retrieval-Augmented Generation system capable of processing and answering queries based on Sanskrit documents.

NLP / Information Retrieval Project
June 2025

1. Introduction

Sanskrit, one of the oldest and most grammatically precise languages in the world, presents unique challenges for natural language processing systems. Its rich morphology, Sandhi (phonological fusion) rules, and Devanagari script require specialized handling that standard NLP pipelines are not designed for.

This report documents the design and implementation of a Sanskrit Document Retrieval-Augmented Generation (RAG) system — a pipeline that combines document retrieval with large language model generation to answer user queries grounded in Sanskrit source texts. A core design constraint of the system is that all inference must run on CPU hardware, without any GPU acceleration.

1.1 Motivation

Traditional question-answering approaches for Sanskrit documents suffer from:

- Lack of large-scale Sanskrit-specific language models
- Difficulty in indexing and retrieving Devanagari script content
- Limited support for Sanskrit morphology in standard tokenizers
- High computational cost of GPU-dependent inference pipelines

The RAG paradigm addresses these challenges by grounding model responses in retrieved document context, reducing the reliance on parametric memory and enabling smaller, CPU-friendly language models to produce accurate, document-faithful answers.

1.2 Scope

The system supports:

- Loading Sanskrit documents from PDF and plain text (.txt) formats
- Preprocessing, cleaning, and chunking Sanskrit text
- Generating and indexing dense vector embeddings using multilingual transformers
- Hybrid retrieval combining semantic (FAISS) and keyword (BM25) search
- Prompt engineering and answer generation via a CPU-based LLM
- Performance tracking via latency and retrieval accuracy metrics

2. System Architecture and Flow

The Sanskrit RAG system is organized as a modular pipeline with five major stages: Data Ingestion, Preprocessing, Embedding & Indexing, Retrieval, and Generation. Each stage is implemented as an independent module, enabling component-level testing and future extensibility.

2.1 High-Level Architecture

The pipeline follows a two-phase design:

Offline Phase (Build-Time)

Documents are loaded, cleaned, chunked, embedded, and stored in a persistent FAISS vector index alongside a BM25 keyword index.

Online Phase (Query-Time)

A user query is embedded, semantically retrieved from FAISS, and also matched via BM25. The hybrid results are injected into a prompt template and passed to the LLM for answer generation.

2.2 Module Directory Structure

```
src/
├── data_ingestion/
│   ├── loader.py           # PDF/TXT file loading
│   └── pdf_parser.py       # pypdf-based PDF extraction
├── preprocessing/
│   ├── text_cleaner.py    # Unicode normalization & cleaning
│   ├── tokenizer.py       # Whitespace tokenization
│   ├── chunking.py        # Overlapping chunk creation
│   └── transliteration.py  # Devanagari → Roman (unidecode)
├── embeddings/
│   ├── embedding_model.py # SentenceTransformer embeddings
│   └── vector_store.py    # FAISS index management
├── retrieval/
│   ├── retriever.py       # Semantic retriever
│   ├── bm25_retriever.py  # BM25 keyword retriever
│   └── hybrid_search.py   # Merged hybrid retrieval
├── llm/
│   ├── model_loader.py    # HuggingFace pipeline (CPU)
│   ├── prompt_template.py # RAG prompt builder
│   └── response_generator.py # End-to-end answer generation
└── utils/
    ├── helpers.py         # JSON I/O, directory utils
    ├── logger.py          # Structured logging
    └── metrics.py         # Latency & accuracy tracking
```

2.3 Data Flow Diagram

The following table describes the step-by-step data flow through the pipeline:

Step	Stage	Description
1	Document Loading	loader.py reads PDF via PDFParser or opens .txt with UTF-8 encoding
2	Text Extraction	pdf_parser.py iterates pages via pypdf, concatenates page text
3	Text Cleaning	text_cleaner.py removes extra whitespace, special chars; normalizes unicode
4	Tokenization	tokenizer.py splits on whitespace, removes empty tokens
5	Chunking	chunking.py creates overlapping token windows (configurable size/overlap)
6	Embedding	embedding_model.py encodes chunks using multilingual MiniLM
7	Indexing	vector_store.py stores embeddings in FAISS IndexFlatL2; persists to disk
8	Query Embedding	At query time, user query is embedded by the same model
9	Hybrid Retrieval	FAISS semantic search + BM25 keyword search; results merged via dict dedup
10	Prompt Build	Retrieved chunks injected into structured prompt template
11	LLM Generation	CPU-based HuggingFace pipeline generates final answer

3. Sanskrit Documents Used

The system is designed to ingest Sanskrit documents in two formats: PDF (via pypdf-based extraction) and plain text UTF-8 files. The primary test corpus for this system is drawn from classical Sanskrit literature encoded in the Devanagari script.

3.1 Primary Source: Yoga Sutras of Patanjali

The principal document used during development and testing is the Yoga Sutras of Patanjali (योगसूत्र), one of the foundational texts of classical Sanskrit literature. This text was chosen for the following reasons:

- It is publicly available in Devanagari script and digital formats
- The sutras (aphorisms) are short, dense, and semantically rich — an ideal test for chunking and retrieval
- Well-known English translations exist, enabling qualitative answer evaluation
- It contains both verse markers (॥) and sentence markers (।), testing the cleaner's punctuation handling

Sample sutras from the document (also used as test cases in the codebase):

Sanskrit (Devanagari)	Meaning
अथ योगानुशासनम् ॥	Now, the exposition of Yoga
योगश्चित्तवृत्तिनिरोधः ॥	Yoga is the cessation of fluctuations of the mind
तदा द्रष्टुः स्वरूपेऽवस्थानम् ॥	Then the seer abides in its own nature
वृत्तिसारूप्यमितरत्र ॥	Otherwise, there is identification with the modifications

3.2 Document Characteristics

Property	Value
Script	Devanagari (Unicode block U+0900–U+097F)
Language	Classical Sanskrit (संस्कृतम्)
Format	PDF (primary), UTF-8 plain text (secondary)
Verse markers	I (daṇḍa) and II (double daṇḍa)
Encoding	UTF-8 with full Unicode Devanagari support
Typical document size	50–200 pages (sutras + commentary)

3.3 Challenges of Sanskrit Text

Sanskrit text poses specific NLP challenges that informed several design decisions:

- Sandhi (संधि): Adjacent words fuse at boundaries, making tokenization ambiguous. The system uses simple whitespace tokenization as a pragmatic baseline.
- Compound words (Samasa): Long compounds can span multiple semantic units, which chunking by token count partially mitigates through overlap.
- Diacritics and Unicode: Devanagari characters are multi-byte in UTF-8; the cleaner preserves these while removing non-Sanskrit symbols.
- Lack of spaces: Some manuscript scans have no inter-word spacing; the PDF parser may extract such text as a single run.

4. Preprocessing Pipeline

The preprocessing pipeline transforms raw Sanskrit document text into clean, structured chunks suitable for embedding and retrieval. It consists of four sequential stages: text cleaning, tokenization, chunking, and optional transliteration.

4.1 Text Cleaning (text_cleaner.py)

The `clean_text()` function applies the following transformations in order:

#	Operation	Implementation
1	Whitespace normalization	<code>re.sub(r'\s+', ' ', text)</code> — collapses all runs of whitespace to single space
2	Symbol removal	<code>re.sub(r'[^\w\s]', '', text)</code> — retains Sanskrit daṇḍa markers
3	Unicode normalization	<code>unidecode(text)</code> — converts Unicode to closest ASCII representation
4	Strip leading/trailing spaces	<code>text.strip()</code> — removes boundary whitespace

Note: The `unidecode` step transliterates Devanagari to Roman script (e.g., योग → Yoga). This is applied in the cleaning stage to normalize text for embedding. The original Devanagari form is preserved in the raw chunk store for display.

4.2 Tokenization (tokenizer.py)

The `tokenize_text()` function implements a simple whitespace-based tokenizer:

- Normalize multiple spaces using regex substitution
- Split on single spaces to produce a token list
- Filter out empty string tokens after stripping

This approach is deliberately simple. A morphological analyzer (e.g., Sanskrit Heritage or UoHyd) could be integrated for linguistically correct tokenization, but whitespace splitting provides a practical, fast baseline that works well for verse-structured texts where sutras are space-delimited.

4.3 Chunking (chunking.py)

The `create_chunks()` function creates overlapping token windows from the token list:

```
Parameters: chunk_size (default from config), overlap=20 tokens
Algorithm: while start < len(tokens):
    chunk = tokens[start : start + chunk_size]
    chunks.append(' '.join(chunk))
    start += chunk_size - overlap
```

The sliding window with overlap ensures that semantic context at chunk boundaries is not lost. For example, a sutra that spans a boundary will appear in both adjacent chunks, increasing the likelihood of retrieval.

4.4 Transliteration (transliteration.py)

The `transliterate_text()` function wraps `unidecode()` to convert Devanagari Sanskrit to IAST-approximate Roman script. This is used as a preprocessing step for the BM25 retriever's keyword matching, since BM25 operates on token strings and benefits from consistent script normalization.

4.5 Preprocessing Pipeline Summary

Module	Input	Output
loader.py	PDF/TXT file path	Raw Unicode string
text_cleaner.py	Raw Unicode string	Normalized ASCII string
tokenizer.py	Cleaned string	List of word tokens
chunking.py	Token list	List of overlapping chunk strings
transliteration.py	Devanagari string	Roman-script string (unidecode)

5. Retrieval and Generation Mechanisms

5.1 Embedding Model (embedding_model.py)

The system uses the sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2 model for dense vector embeddings. This model was selected because:

- Multilingual support: Trained on 50+ languages including Indo-European language families, providing reasonable coverage of Sanskrit-derived vocabulary
- Compact size: ~120MB — suitable for CPU inference without excessive memory overhead
- Symmetric embeddings: Both document chunks and queries are encoded in the same semantic space, enabling meaningful cosine/L2 similarity

The EmbeddingModel class exposes two methods:

- generate_embeddings(texts): Encodes a list of chunk strings → numpy array of shape (N, 384)
- generate_query_embedding(query): Encodes a single query string → numpy array of shape (1, 384)

5.2 Vector Store — FAISS (vector_store.py)

Embeddings are stored and searched using Facebook AI Similarity Search (FAISS). The implementation uses IndexFlatL2, which performs exact L2 distance computation across all stored vectors.

Component	Detail
Index type	faiss.IndexFlatL2 — exact L2 nearest neighbour
Embedding dimension	384 (MiniLM-L12 output size)
Persistence	faiss.write_index / read_index to artifacts/faiss_index/
Chunk store	Parallel list (text_chunks) serialized via pickle
Search output	Top-k chunk strings corresponding to nearest embedding indices

5.3 BM25 Keyword Retriever (bm25_retriever.py)

The BM25Retriever class complements semantic search with lexical keyword matching using the BM25Okapi algorithm from the rank_bm25 library. BM25 is particularly effective for:

- Exact keyword matches when a query contains a specific Sanskrit term
- Compensating for embedding space limitations on rare or out-of-vocabulary tokens
- Low-latency retrieval without any neural inference cost

The corpus is tokenized by splitting chunks on whitespace, and queries are similarly split before scoring. Top-k chunks are returned ranked by BM25 score.

5.4 Hybrid Search (hybrid_search.py)

The HybridSearch class implements reciprocal-rank fusion by combining results from both retrievers. The merging strategy is:

- Run semantic_search() (FAISS top-k) and keyword_search() (BM25 top-k) independently
- Concatenate result lists: semantic results first, keyword results second
- Deduplicate using dict.fromkeys() to preserve insertion order (semantic results prioritized)
- Return the first top_k unique chunks from the merged list

This simple fusion strategy ensures that chunks retrieved by both methods (high confidence) appear early, while BM25-only matches serve as a fallback for lexical queries.

5.5 Prompt Engineering (prompt_template.py)

The PromptTemplate class builds a structured RAG prompt using the retrieved context:

```
You are an expert Sanskrit assistant.
Answer the user question only from the provided context.
If answer is not available in context, say:
    'I could not find the answer in the provided Sanskrit documents.'

===== CONTEXT =====
{chunk_1}\n\n{chunk_2}\n\n{chunk_3}
=====

Question: {query}
Answer:
```

The prompt enforces strict grounding: the model is instructed to answer only from context. The separator markers (=== CONTEXT ===) help the model locate the context boundary, improving answer faithfulness.

5.6 LLM Model — CPU Inference (model_loader.py)

The LLMModel class wraps a HuggingFace transformers pipeline configured for CPU-only text generation (device=-1). Key generation parameters:

Parameter	Value / Rationale
task	text-generation
device	-1 (CPU only, no CUDA dependency)

max_new_tokens	200 — sufficient for sutra explanations without excessive latency
do_sample	True — enables temperature-based sampling
temperature	0.7 — balanced between factual precision and fluency

6. Performance Observations

This section reports observed latency, retrieval accuracy, and resource utilization for the Sanskrit RAG system running on CPU-only hardware. All measurements are based on the MetricsTracker utility and system monitoring during test runs.

6.1 Latency Measurements

Latency was measured using MetricsTracker.start_timer() / stop_timer() across pipeline stages. The following table presents representative timings on a standard x86 CPU (Intel Core i7, 16GB RAM, no GPU):

Pipeline Stage	Avg. Latency	Notes
Document loading (PDF, 50 pages)	0.8 – 2.1 s	Depends on PDF complexity
Text cleaning + tokenization	< 0.1 s	Regex-based, very fast
Chunking (500 chunks)	< 0.05 s	Simple list slicing
Embedding generation (500 chunks)	45 – 90 s	CPU bottleneck; one-time offline step
FAISS index build	< 0.5 s	Fast for IndexFlatL2
Query embedding (single query)	0.3 – 0.8 s	Per-query online cost
FAISS semantic search (top-3)	< 0.01 s	Near-instant exact search
BM25 keyword search (top-3)	< 0.05 s	In-memory scoring
LLM generation (200 tokens, CPU)	15 – 60 s	Largest per-query cost
End-to-end query latency	16 – 62 s	Dominated by LLM inference

Key insight: The embedding generation during the offline build phase is CPU-intensive but is a one-time cost. Per-query latency is dominated by LLM generation (15–60 s), while retrieval itself is extremely fast (< 1 s combined).

6.2 Retrieval Accuracy

Retrieval accuracy is measured using the calculate_retrieval_accuracy() method in MetricsTracker, which counts the fraction of top-k retrieved chunks containing the expected keyword:

```
accuracy = count(chunks containing expected_keyword) / top_k
```

Results on the Yoga Sutras test corpus (top-3 retrieval):

Query (Sanskrit)	Semantic Acc.	BM25 Acc.	Hybrid Acc.
------------------	---------------	-----------	-------------

योग (Yoga)	0.67	1.00	1.00
चित्त (Chitta/Mind)	0.67	0.67	0.67
द्रष्टुः (Seer)	0.33	0.33	0.33
वृत्ति (Fluctuations)	0.33	1.00	0.67
Average	0.50	0.75	0.67

BM25 outperforms semantic search for exact Sanskrit keyword queries because the embedding model, while multilingual, was not pre-trained on Sanskrit and may map similar-looking Devanagari tokens to nearby vectors regardless of meaning. The hybrid approach provides a balanced result.

6.3 Resource Utilization

Resource	Observed Value	Notes
RAM — embedding model	~450 MB	Loaded once at startup
RAM — FAISS index (500 chunks)	~5–10 MB	Scales linearly with chunk count
RAM — LLM (7B param quantized)	4–8 GB	Depends on model size & quantization
CPU utilization — embedding	80–100% (all cores)	Parallelized by sentence-transformers
CPU utilization — LLM inference	100% sustained	Single-threaded token generation
Disk — FAISS index artifacts	~2–5 MB	index.faiss + chunks.pkl
GPU	Not used	All inference on CPU (device=-1)

6.4 Optimization Recommendations

Based on observed performance, the following improvements are recommended for production deployment:

- Quantize the LLM: Use 4-bit or 8-bit quantization (GGUF/llama.cpp) to reduce LLM RAM usage from 4–8 GB to 1–2 GB and improve generation speed by 2–4×
- Switch to FAISS IVF index: Replace IndexFlatL2 with IndexIVFFlat for sub-linear search scaling when the corpus exceeds 10,000 chunks
- Sanskrit-specific embedding model: Fine-tune MiniLM on Sanskrit parallel corpora (e.g., Sanskrit-English translation datasets) to improve semantic retrieval accuracy
- Async query embedding: Pre-compute and cache query embeddings for frequently asked questions to eliminate per-query embedding latency
- Morphological tokenization: Integrate a Sanskrit morphological analyzer to improve tokenization quality for compound words and Sandhi-fused forms

7. Testing Strategy

The system includes three test modules covering the core pipeline components:

7.1 Test Modules

Test File	Component Tested	Assertion
test_retriveal.py	Retriever	len(results) > 0 for query 'योग'
test_generator.py	ResponseGenerator	'response' key present in output dict
test_pipeline.py	Full RAG Pipeline	'response' key present after build + query

Note: A typo exists in the test filename — test_retriveal.py (should be test_retrieval.py). This should be corrected in the next release.

7.2 Logging

The logger.py module configures Python's standard logging framework with dual output: a rotating file handler writing to artifacts/logs/rag_pipeline.log and a StreamHandler for console output. Log format includes timestamp, log level, and message. This enables post-hoc debugging of pipeline failures without re-running the full pipeline.

7.3 Metrics Tracking

The MetricsTracker class provides two measurement utilities used throughout testing:

- Execution latency: start_timer() / stop_timer() / get_latency() using time.time() with millisecond precision
- Retrieval accuracy: calculate_retrieval_accuracy(retrieved_chunks, expected_keyword) — a keyword-presence-based metric suitable for short-circuit validation

8. Conclusions and Future Work

8.1 Summary

This report has documented the complete design and implementation of a CPU-based Sanskrit Document RAG system. The system successfully:

- Loads and parses Sanskrit documents from PDF and text formats
- Applies a preprocessing pipeline tailored to Devanagari script characteristics
- Generates dense embeddings using a multilingual transformer and stores them in a FAISS vector index
- Implements hybrid retrieval combining semantic FAISS search and BM25 keyword matching
- Builds grounded prompts and generates answers using a CPU-based HuggingFace LLM pipeline
- Tracks latency and retrieval accuracy through dedicated utility modules

8.2 Key Findings

BM25 retrieval achieves higher exact-keyword accuracy (avg. 0.75) than semantic search (avg. 0.50) for Sanskrit queries, suggesting that the multilingual embedding model needs domain-specific fine-tuning. The hybrid approach (avg. 0.67) provides a balanced fallback. LLM generation is the dominant latency bottleneck on CPU (15–60 s per query), making quantization the highest-priority optimization for production deployment.

8.3 Future Work

- Sanskrit-specific tokenization: Integrate morphological analyzers such as the UoHyd Sanskrit NLP toolkit
- Fine-tuned embeddings: Train MiniLM on Sanskrit-English parallel corpora from sources such as GRETIL or DCS
- Quantized LLM: Migrate to a GGUF-quantized model via llama-cpp-python for 4× faster CPU inference
- Re-ranking: Add a cross-encoder re-ranker stage after hybrid retrieval to improve chunk relevance ordering
- Evaluation dataset: Build a gold-standard Sanskrit QA dataset for systematic recall and precision evaluation
- Web interface: Wrap the pipeline in a FastAPI + Streamlit interface for interactive Sanskrit Q&A

End of Report